

## EXPERIMENT NO. 10

### Character, Digit, or Face Classification using Convolutional

### Neural Networks (CNNs)

NAME : Abhigyan Dubey

CUID : CU24250059

COURSE : BTECH CSE

SECTION : B

ROLL NO. : 02

### AIM

To design, train, and evaluate a Convolutional Neural Network

(CNN) for handwritten digit classification using the MNIST

dataset and analyze its performance using suitable metrics.

STEP 1: IMPORT REQUIRED LIBRARIES

```
import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf

from tensorflow.keras import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from sklearn.metrics import confusion_matrix, classification_report
import seaborn as sns
```

STEP 2: LOAD THE MNIST DATASET

```
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.mnist.load_data()

print("Training images:", x_train.shape)
print("Testing images :", x_test.shape)
```

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-
datasets/mnist.npz
11490434/11490434 [0m 0s 32m 0s] 0m 0s/step
Training images: (60000, 28, 28)
Testing images : (10000, 28, 28)

```

### STEP 3: PREPROCESS THE DATA

```

# Normalize pixel values from 0-255 to 0-1
x_train = x_train.astype("float32") / 255.0
x_test = x_test.astype("float32") / 255.0

# Add channel dimension for CNN
x_train = x_train.reshape(-1, 28, 28, 1)
x_test = x_test.reshape(-1, 28, 28, 1)

# Split training data into training and validation data
x_val = x_train[-10000:]
y_val = y_train[-10000:]

x_train = x_train[:-10000]
y_train = y_train[:-10000]

print("Training data after split:", x_train.shape)
print("Validation data:", x_val.shape)
print("Testing data:", x_test.shape)

```

```

Training data after split: (50000, 28, 28, 1)
Validation data: (10000, 28, 28, 1)
Testing data: (10000, 28, 28, 1)

```

### STEP 4: DISPLAY SAMPLE IMAGES


```

plt.figure(figsize=(8, 6))

for i in range(9):
    plt.subplot(3, 3, i + 1)
    plt.imshow(x_train[i].reshape(28, 28), cmap="gray")
    plt.title("Digit: " + str(y_train[i]))
    plt.axis("off")

```

```
plt.tight_layout()  
plt.show()
```



## STEP 5: DESIGN THE CNN MODEL

```
model = Sequential([  
    Conv2D(32, (3, 3), activation="relu", input_shape=(28, 28, 1)),  
    MaxPooling2D((2, 2)),  
  
    Conv2D(64, (3, 3), activation="relu"),  
    MaxPooling2D((2, 2)),  
  
    Flatten(),
```

```

Dense(128, activation="relu"),
Dropout(0.3),

Dense(10, activation="softmax")
])
# Display model architecture
model.summary()

```

```

/usr/local/lib/python3.13/dist-
packages/keras/src/layers/convolutional/base_conv.py:113: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer
using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

**Model: "sequential"**

| Layer (type)                   | Output Shape       | Param # |
|--------------------------------|--------------------|---------|
| conv2d (Conv2D)                | (None, 26, 26, 32) | 320     |
| max_pooling2d (MaxPooling2D)   | (None, 13, 13, 32) | 0       |
| conv2d_1 (Conv2D)              | (None, 11, 11, 64) | 18,496  |
| max_pooling2d_1 (MaxPooling2D) | (None, 5, 5, 64)   | 0       |
| flatten (Flatten)              | (None, 1600)       | 0       |
| dense (Dense)                  | (None, 128)        | 204,928 |
| dropout (Dropout)              | (None, 128)        | 0       |
| dense_1 (Dense)                | (None, 10)         | 1,290   |

**Total params: 225,034 (879.04 KB)**

**Trainable params:** 225,034 (879.04 KB)

**Non-trainable params:** 0 (0.00 B)

## STEP 6: COMPILE THE MODEL

```
model.compile(
    optimizer="adam",
    loss="sparse_categorical_crossentropy",
    metrics=["accuracy"]
)
```

## STEP 7: TRAIN THE CNN MODEL

```
history = model.fit(
    x_train,
    y_train,
    validation_data=(x_val, y_val),
    epochs=5,
    batch_size=128
)
```

Epoch 1/5

1/391 [32m] 1/37m 1/40s 98ms/step - accuracy: 0.9148 - loss: 0.2800 - val\_accuracy: 0.9777 - val\_loss: 0.0742

Epoch 2/5

1/391 [32m] 1/37m 1/38s 98ms/step - accuracy: 0.9747 - loss: 0.0797 - val\_accuracy: 0.9846 - val\_loss: 0.0559

Epoch 3/5

1/391 [32m] 1/37m 1/40s 96ms/step - accuracy: 0.9824 - loss: 0.0572 - val\_accuracy: 0.9879 - val\_loss: 0.0462

Epoch 4/5

1/391 [32m] 1/37m 1/38s 96ms/step - accuracy: 0.9869 - loss: 0.0438 - val\_accuracy: 0.9887 - val\_loss: 0.0385

Epoch 5/5

```
1m391/391[0m 32m[0m 37m[0m 1m42s[0m 106ms/step -  
accuracy: 0.9889 - loss: 0.0359 - val_accuracy: 0.9906 - val_loss: 0.0344
```

## STEP 8: DISPLAY TRAINING AND VALIDATION ACCURACY

```
plt.figure(figsize=(8, 5))  
  
plt.plot(history.history["accuracy"], label="Training Accuracy")  
plt.plot(history.history["val_accuracy"], label="Validation Accuracy")  
  
plt.title("Training and Validation Accuracy")  
plt.xlabel("Epoch")  
plt.ylabel("Accuracy")  
plt.legend()  
plt.grid()  
  
plt.show()
```




## STEP 9: DISPLAY TRAINING AND VALIDATION LOSS

```
plt.figure(figsize=(8, 5))

plt.plot(history.history["loss"], label="Training Loss")
plt.plot(history.history["val_loss"], label="Validation Loss")

plt.title("Training and Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid()

plt.show()
```



## STEP 10: EVALUATE MODEL ON TEST DATA

```
test_loss, test_accuracy = model.evaluate(
    x_test,
    y_test,
    verbose=0
```

```
)  
  
print("\nTest Accuracy:", test_accuracy)  
print("Test Loss:", test_loss)
```

```
Test Accuracy: 0.9908000230789185  
Test Loss: 0.027723412960767746
```

## STEP 11: MAKE PREDICTIONS

```
predictions = model.predict(x_test, verbose=0)  
  
y_pred = np.argmax(predictions, axis=1)
```

## STEP 12: CALCULATE PRECISION, RECALL AND F1-SCORE

```
report = classification_report(  
    y_test,  
    y_pred,  
    digits=4  
)
```

```
print("\nClassification Report:")  
print(report)
```

### Classification Report:

|   | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.9899    | 0.9959 | 0.9929   | 980     |
| 1 | 0.9921    | 0.9938 | 0.9930   | 1135    |
| 2 | 0.9904    | 0.9952 | 0.9928   | 1032    |
| 3 | 0.9901    | 0.9931 | 0.9916   | 1010    |
| 4 | 0.9949    | 0.9919 | 0.9934   | 982     |
| 5 | 0.9844    | 0.9877 | 0.9860   | 892     |
| 6 | 0.9937    | 0.9885 | 0.9911   | 958     |
| 7 | 0.9922    | 0.9854 | 0.9888   | 1028    |
| 8 | 0.9938    | 0.9908 | 0.9923   | 974     |
| 9 | 0.9861    | 0.9851 | 0.9856   | 1009    |

|              |        |        |        |       |
|--------------|--------|--------|--------|-------|
| accuracy     |        |        | 0.9908 | 10000 |
| macro avg    | 0.9907 | 0.9907 | 0.9907 | 10000 |
| weighted avg | 0.9908 | 0.9908 | 0.9908 | 10000 |

### STEP 13: CONFUSION MATRIX


```
cm = confusion_matrix(y_test, y_pred)
```

```
plt.figure(figsize=(8, 6))
```

```
sns.heatmap(  
    cm,  
    annot=True,  
    fmt="d",  
    cmap="Blues"  
)
```

```
plt.title("Confusion Matrix")  
plt.xlabel("Predicted Label")  
plt.ylabel("Actual Label")
```

```
plt.show()
```



#### STEP 14: PREDICTION ON UNSEEN IMAGES

```
plt.figure(figsize=(10, 8))

for i in range(12):
    plt.subplot(3, 4, i + 1)

    plt.imshow(
        x_test[i].reshape(28, 28),
        cmap="gray"
    )


    predicted_digit = y_pred[i]
    actual_digit = y_test[i]

    plt.title(
        "Actual: " + str(actual_digit) +
        "\nPredicted: " + str(predicted_digit)
```

```
)

plt.axis("off")

plt.tight_layout()
plt.show()
```



## STEP 15: FIND MISCLASSIFIED IMAGES

```
wrong_indices = np.where(y_pred != y_test)[0]

print("\nNumber of misclassified images:", len(wrong_indices))

plt.figure(figsize=(10, 8))

for i in range(min(12, len(wrong_indices))):
```

```
index = wrong_indices[i]

plt.subplot(3, 4, i + 1)


plt.imshow(
    x_test[index].reshape(28, 28),
    cmap="gray"
)

plt.title(
    "Actual: " + str(y_test[index]) +
    "\nPredicted: " + str(y_pred[index])
)

plt.axis("off")

plt.tight_layout()
plt.show()
```

Number of misclassified images: 92



## QUESTIONS AND ANSWERS

Q1. What is image classification? How does it differ from object detection and image segmentation?

Answer:

Image classification assigns one or more class labels to an image.

Object detection identifies objects and also provides their locations using bounding boxes. Image segmentation assigns a class to individual pixels or regions of an image.

Q2. Explain the architecture and working principle of a CNN.

Answer:

A CNN normally contains convolutional layers, activation functions,

pooling layers, flattening, fully connected layers, and an output layer. Convolutional layers extract image features, pooling reduces feature dimensions, and fully connected layers perform classification.

Q3. What is the role of Convolutional Layers, Pooling Layers, and Fully Connected Layers in a CNN?

Answer:

Convolutional layers detect features such as edges and shapes. Pooling layers reduce the size of feature maps and retain important information. Fully connected layers use the extracted features to perform the final classification.

Q4. Why is image normalization performed before training a deep learning model?

Answer:

Image normalization scales pixel values to a smaller range such as 0 to 1. It helps the neural network train faster and makes the optimization process more stable.

Q5. Explain the purpose of activation functions such as ReLU and Softmax in CNNs.

Answer:

ReLU introduces non-linearity into the network and is commonly used in hidden layers. Softmax converts the final output values into probabilities for the different classes.

Q6. What is a Confusion Matrix? How is it used to evaluate classification performance?

Answer:

A confusion matrix is a table showing actual classes against predicted classes. It shows correct predictions and different types of misclassification for each class.

Q7. Differentiate between Accuracy, Precision, Recall, and F1-Score.

Answer:

Accuracy is the proportion of all predictions that are correct. Precision measures how many predicted positive samples are actually positive. Recall measures how many actual positive samples are detected. F1-Score is the harmonic mean of precision and recall.

Q8. Compare traditional feature-based image classification methods such as SIFT/HOG + SVM with CNN-based classification.

Answer:

SIFT/HOG + SVM requires manually designed or extracted features before classification. CNNs automatically learn useful features directly from training images. CNNs generally require more computational resources and training data but can learn complex features automatically.

Q9. Mention five real-world applications of CNN-based image classification.

Answer:

1. Handwritten digit recognition
2. Facial recognition
3. Medical image analysis
4. Biometric authentication
5. Intelligent document processing

Q10. How can Data Augmentation, Transfer Learning, and Hyperparameter Tuning improve deep learning image classification?

Answer:

Data augmentation creates variations of training images and can reduce overfitting. Transfer learning uses knowledge learned from an existing model and can reduce training requirements. Hyperparameter tuning helps select suitable learning rate, batch size, number of layers, epochs, and other training parameters.